

Jiwon Jeong

AI Research Engineer

✉ jwjw9603@gmail.com |  github.com/jw9603 |  [LinkedIn](#) |  [Hugging Face](#) |  jw9603.github.io

RESEARCH INTERESTS

Attention Architecture and Efficiency, Mechanistic Interpretability, LLM Agent Evaluation

EDUCATION

Sungkyunkwan University (SKKU)

M.S. in Artificial Intelligence

Suwon, South Korea

Mar. 2022 – Aug. 2024

- Advisor: Prof. Hogun Park
- Research: Commonsense QA, Logical Fallacy Detection

Kookmin University

B.S. in Electronic Engineering

Seoul, South Korea

Mar. 2016 – Feb. 2022

PUBLICATIONS

Under Review

2026

- Two papers under double-blind review at ICLR 2027 (titles withheld)

NeurIPS 2026 Workshop on SLMs for Agentic Systems (SLM-Agents) | *Poster*

2026

- Beyond Query Linearity: The Case for Nonlinear Queries
- **Jiwon Jeong**, Dun Li Chan, Andrew Ceniccola, Rahul Sharma, Simon Sang, Parva Mehta, Pietro Di Nicco, Karim Saraipour, Marko Karbevski

The 7th Korea Artificial Intelligence Conference

2026

- Beyond Recovery Rate: Post-Recovery Verification and Reporting Gaps in Failure-Aware Evaluation of LLM Agents
- **Jiwon Jeong**, Hyunmook Cha
- [PDF](#)

NAACL 2025 Findings

2025

- Large Language Models Are Better Logical Fallacy Reasoners with Counterargument, Explanation, and Goal-Aware Prompt Formulation
- **Jiwon Jeong**, Hyeju Jang, Hogun Park
- [arXiv](#)

KSC 2022 | *Outstanding Paper Award (Top 10%)*

2022

- Improving Commonsense-based QA Model through a Cycle-Encoder
- **Jiwon Jeong**, Soyoung Lee, Hogun Park
- [DBpia](#)

PROFESSIONAL EXPERIENCE

AI Research Engineer

Lotte Innovate – AI Tech LAB, Language AI Team

Sep. 2025 – Present

Seoul, South Korea

- AI Research Engineer responsible for LLM research and enterprise deployment

TEACHING & MENTORING

Online Mentor

Feb. 2026

LIKELION – AI NLP Engineer Intensive Course, 3rd Cohort

Remote

- Mentored 4 teams on AI-based service development projects using corporate and public data
- Provided code review, project coaching, and Q&A support via Discord

Teaching Assistant

2023 – 2024

Sungkyunkwan University

Suwon, South Korea

- TA for Language Model course (SWE3032-41) and Deep Learning advanced courses for non-CS educators

ACTIVITIES

- Runner – “Reading LLM Technical Reports” (13th Cohort)** Oct. 2026 – Present
Pseudo-Lab ([link](#)) – Non-profit AI research community *Remote*
- 16-week study reading recent LLM technical reports in depth: architecture, data, training, and whether each claim is supported by its evidence
 - The cohort plans to write a survey paper on LLM technical reports after the study
- Research Contributor – ConspiracyBench** Sep. 2026 – Present
Stationhouse ([link](#)) – Open research community for open model research *Remote*
- Contributing to schema design and benchmark item submission for an LLM evaluation benchmark
- Runner – “Deep Dive into LLM”** Mar. 2026 – Jul. 2026
Pseudo-Lab ([link](#)) – Non-profit AI research community *Remote*
- Studied Stanford CS336: Language Modeling from Scratch (2026) through weekly online discussions
- AI Researcher** Aug. 2025 – Jan. 2026
B-Peach LAB (Tech for Impact) *Remote*
- Designed evaluation metrics for AI-powered text adaptation service improving information accessibility for slow learners
 - Conducted research on adaptation quality assessment

PROJECTS

- Open Source Contribution: Hugging Face Transformers** | *Python, Transformers* Mar. 2026 – Present
- Removed six dead beam-search dummy classes left over from v5 cleanup, fixing a repository consistency check failure on PyTorch-less installs ([PR #45722](#), merged Apr. 2026)
 - Reported and fixed a protobuf decode error masking bug in tokenizer loading ([PR #45486](#), merged Apr. 2026)
 - Fixed a bug in continuous batching for multimodal models. Added input preprocessing and fallback mechanism for non-text-only models ([PR #44436](#), merged Mar. 2026)
- MCP Server for ChatGPT Apps** | *Python, FastAPI, MCP* Jan. 2026 – Present
- Built MCP (Model Context Protocol) servers for ChatGPT Apps (OpenAI)
 - Designed middleware architecture connecting frontend, backend, and LLM-powered enterprise solutions
- Korean LLM Benchmark** | *Python, LLM, Evaluation* Dec. 2025 – Jan. 2026
- Evaluation framework for Korean LLM performance with standardized metrics and reproducible pipelines

AWARDS

- KSC 2022 Outstanding Paper Award** | *Top 10% of accepted papers* 2022
- Text Ethics Hackathon – 2nd Place** Nov. 2021 – Dec. 2021
- Developed ELECTRA-based classification model for detecting hate speech and biased text

TECHNICAL SKILLS

Languages: Python
Frameworks: PyTorch, Hugging Face Transformers, FastAPI, LangChain, Docker
Research: NLP, Transformer Architecture, Mechanistic Interpretability, LLM Fine-tuning, Prompt Engineering, Model Evaluation, MCP
Tools: Git, Linux, Weights & Biases, Jupyter