

Jiwon Jeong

AI Research Engineer

✉ jwjw9603@gmail.com | github.com/jw9603 | [LinkedIn](https://www.linkedin.com/in/jw9603) | [Hugging Face](https://huggingface.co/jw9603) | jw9603.github.io

RESEARCH INTERESTS

Attention Architecture and Efficiency, Mechanistic Interpretability, Text Simplification and Intralingual Translation for Slow Learners, LLM Agent Evaluation

EDUCATION

Sungkyunkwan University (SKKU)

M.S. in Artificial Intelligence

Suwon, South Korea

Mar. 2022 – Aug. 2024

- Advisor: Prof. Hogun Park
- Research: Commonsense QA, Logical Fallacy Detection

Kookmin University

B.S. in Electronic Engineering

Seoul, South Korea

Mar. 2016 – Feb. 2022

PUBLICATIONS

Under Review

2026

- Two papers under double-blind review at ICLR 2027 (titles withheld)

NeurIPS 2026 Workshop on SLMs for Agentic Systems (SLM-Agents) | *Poster*

2026

- Beyond Query Linearity: The Case for Nonlinear Queries
- **Jiwon Jeong**, Dun Li Chan, Andrew Ceniccola, Rahul Sharma, Simon Sang, Parva Mehta, Pietro Di Nicco, Karim Saraipour, Marko Karbevski

The 7th Korea Artificial Intelligence Conference

2026

- Beyond Recovery Rate: Post-Recovery Verification and Reporting Gaps in Failure-Aware Evaluation of LLM Agents
- **Jiwon Jeong**, Hyunmook Cha
- [PDF](#)

NAACL 2025 Findings

2025

- Large Language Models Are Better Logical Fallacy Reasoners with Counterargument, Explanation, and Goal-Aware Prompt Formulation
- **Jiwon Jeong**, Hyeju Jang, Hogun Park
- [arXiv](#)

KSC 2022 | *Outstanding Paper Award (Top 10%)*

2022

- Improving Commonsense-based QA Model through a Cycle-Encoder
- **Jiwon Jeong**, Soyoung Lee, Hogun Park
- [DBpia](#)

PROFESSIONAL EXPERIENCE

AI Research Engineer

Lotte Innovate – AI Tech LAB, Language AI Team

Sep. 2025 – Present

Seoul, South Korea

- AI Research Engineer responsible for LLM research and enterprise deployment

TEACHING & MENTORING

Online Mentor

Feb. 2026

LIKELION – AI NLP Engineer Intensive Course, 3rd Cohort

Remote

- Mentored 4 teams on AI-based service development projects using corporate and public data
- Provided code review, project coaching, and Q&A support via Discord

Teaching Assistant

2023 – 2024

Sungkyunkwan University

Suwon, South Korea

- TA for Language Model course (SWE3032-41) and Deep Learning advanced courses for non-CS educators

ACTIVITIES

Research Contributor – ConspiracyBench

Sep. 2026 – Present

Stationhouse ([link](#)) – Open research community for open model research

Remote

- Contributing to schema design and benchmark item submission for an LLM evaluation benchmark

Runner – “Deep Dive into LLM”

Mar. 2026 – Jul. 2026

Pseudo-Lab ([link](#)) – Non-profit AI research community

Remote

- Studied Stanford CS336: Language Modeling from Scratch (2026) through weekly online discussions

AI Researcher

Aug. 2025 – Jan. 2026

B-Peach LAB (Tech for Impact)

Remote

- Designed evaluation metrics for AI-powered text adaptation service improving information accessibility for slow learners
- Conducted research on adaptation quality assessment

PROJECTS

Open Source Contribution: Hugging Face Transformers | Python, Transformers

Mar. 2026 – Present

- Removed six dead beam-search dummy classes left over from v5 cleanup, fixing a repository consistency check failure on PyTorch-less installs ([PR #45722](#), merged Apr. 2026)
- Reported and fixed a protobuf decode error masking bug in tokenizer loading ([PR #45486](#), merged Apr. 2026)
- Fixed a bug in continuous batching for multimodal models. Added input preprocessing and fallback mechanism for non-text-only models ([PR #44436](#), merged Mar. 2026)

MCP Server for ChatGPT Apps | Python, FastAPI, MCP

Jan. 2026 – Present

- Built MCP (Model Context Protocol) servers for ChatGPT Apps (OpenAI)
- Designed middleware architecture connecting frontend, backend, and LLM-powered enterprise solutions

Korean LLM Benchmark | Python, LLM, Evaluation

Dec. 2025 – Jan. 2026

- Evaluation framework for Korean LLM performance with standardized metrics and reproducible pipelines

AWARDS

KSC 2022 Outstanding Paper Award | Top 10% of accepted papers

2022

Text Ethics Hackathon – 2nd Place

Nov. 2021 – Dec. 2021

- Developed ELECTRA-based classification model for detecting hate speech and biased text

TECHNICAL SKILLS

Languages: Python

Frameworks: PyTorch, Hugging Face Transformers, FastAPI, LangChain, Docker

Research: NLP, Transformer Architecture, Mechanistic Interpretability, LLM Fine-tuning, Prompt Engineering, Model Evaluation, MCP

Tools: Git, Linux, Weights & Biases, Jupyter